

Who bears the AI shock? Displacement, poverty and the public finances in the UK*

Vahid Ahmadi

July 2026 — preliminary draft

Abstract

How much an AI labour shock costs the United Kingdom depends on who bears it. I build the first open-source, full tax-benefit microsimulation of AI displacement incidence for the UK: occupation-level AI exposure is attached to Family Resources Survey workers, shocks are simulated across 66 scenarios, and counterfactual incomes are evaluated with PolicyEngine UK, treating who loses work — not just how many — as a scenario axis. In the central scenario, the Exchequer cost is £18.4 billion a year and before-housing-costs poverty rises by 1.87 percentage points. Holding the aggregate shock fixed, the fiscal cost varies from £14.2 billion (uniform incidence) to £24.5 billion (expertise compression); a family shaped by measured UK firm-level displacement patterns costs £32.5 billion. Who bears the shock decides whether the UK gets a fiscal problem or a poverty problem; either way the Gini rises, in all 66 cells and all incidence families.

Keywords: artificial intelligence, labour displacement, microsimulation, poverty, inequality, public finances

JEL classification: C63, D31, H24, H68, J23, I32

1 Introduction

If artificial intelligence displaces a substantial share of UK employment over the next decade, who pays—the displaced workers, the poorest households, or the Exchequer? The answer turns out to depend less on how large the shock is than on a question that policy analysis has so far left implicit: *who bears it*. The UK’s services-intensive economy concentrates employment in the professional, financial, administrative and managerial occupations that score highest on standard measures of AI exposure (Felten et al., 2021; Eloundou et al., 2024; Pizzinelli et al., 2023), while its safety net—means-tested, tapered, and financed overwhelmingly by taxes on labour income—was built to insure shocks that strike the bottom of the income distribution. A displacement shock aimed at the top of the earnings distribution implicates the welfare state twice over: its insurance function faces claimants it rarely serves, and its revenue base sits precisely in the labour income the shock destroys.

Yet the top-loaded incidence implied by exposure indices is an assumption, not a fact. The empirical literature on AI’s labour-market effects has not settled where displacement lands.

*Draft — not for citation.

Exposure-based measures place the risk on high-paid cognitive work (Felten et al., 2021; Eloundou et al., 2024); a competing view holds that AI compresses the returns to accumulated expertise, threatening the wage premia of experienced professionals while opening high-value work to less credentialed workers (Autor, 2024); and the earliest firm-level evidence finds job losses concentrated among *junior* workers within exposed occupations (Klein Teeselink, 2025; Hosseini and Lichtinger, 2026; Brynjolfsson et al., 2025a). These are not refinements of a single estimate. They are different answers to the incidence question, and they imply different claimants at the benefit office and different holes in the tax base.

This paper makes incidence a first-class scenario axis. I build, to my knowledge, the first open-source, full tax-benefit microsimulation of AI labour-shock incidence for the United Kingdom, and the first study I am aware of to run competing incidence assumptions—not just competing shock sizes—through a complete tax and transfer system. Using PolicyEngine UK, an open-source microsimulation model running on the Family Resources Survey 2024–25 with AI exposure assigned to SOC2020 occupations (Felten et al., 2021; Tolan et al., 2021; Eloundou et al., 2024), I simulate a common aggregate shock—seven per cent of employees displaced, a calibration convention within the range implied by Briggs and Kodnani (2023)—under four incidence families: displacement proportional to occupational exposure, displacement concentrated on junior workers, a stylised expertise-compression scenario in which senior premia bear the shock, and a uniform comparator. Around the central calibration I simulate a grid of 66 shock-size combinations crossing displacement rates with wage uplifts, with the capital channel applied throughout. Every result can be regenerated from public code; most comparable exercises rest on closed models and restricted microdata. I then take the framework to policy: a fifth incidence family calibrated to measured UK firm-level displacement patterns, three simulated fiscal responses in the shocked world, a tax-composition accounting of how far corporate profits could refill the revenue hole, and a mapping of the shock to the 650 Westminster constituencies.

Two findings organise the paper. First, one result is robust to everything I vary: the Gini coefficient of equivalised household disposable income rises in all 66 shock-size cells (by 0.06 to 1.55 percentage points) and in all five incidence families (by 0.92 to 1.14 points at the central shock). Inequality rises regardless of who bears the shock—even though, in the exposure family, the market shock itself is progressive in the sense of falling hardest on the highest deciles. The tax-benefit system claws back part of high earners’ losses through forgone tax rather than replacing them through benefits, so a top-loaded market shock still widens the disposable-income distribution.

Second, incidence determines the *form* the cost takes. The same 1.6 million displaced workers cost the Exchequer £18.4 billion when allocated by exposure, £19.4 billion under junior-concentrated incidence, £24.5 billion under expertise compression, and £14.2 billion under uniform incidence—a spread of £10 billion a year on an identical aggregate shock. Poverty moves less: before-housing-costs poverty rises by 1.8 to 1.9 percentage points in three of the four families, but by 2.1 points under expertise compression, which is thus both the most fiscally expensive and the most poverty-raising of the stylised families—and the measured family exceeds it on both margins (Section 4.5). Top-loaded incidence gives the UK a fiscal problem; flatter or junior-concentrated incidence is cheaper for the Treasury but leaves poverty about as high; compression delivers both problems at once. Which world the UK is in cannot yet be settled empirically—so I quantify what is at stake in the disagreement.

The incidence axis matters precisely because exposure indices cannot resolve it. Occupational exposure is age-blind: it cannot distinguish the junior analyst from the senior partner doing

related work, so any scenario allocated by exposure alone spreads displacement across the age structure of exposed occupations, which in the UK peaks among prime-age workers. The junior-concentrated pattern in UK firm data (Klein Teeselink, 2025), US resume histories (Hosseini and Lichtinger, 2026) and US payroll records (Brynjolfsson et al., 2025a) must therefore be imposed on the framework rather than derived from it, and the expertise-compression scenario is an author-designed stress test of the mechanism articulated by Autor (2024) rather than an estimate. I say so plainly throughout: the incidence gradient is an assumption in every family, and the paper’s contribution is the tax-benefit system’s mediation of each.

The remainder of the paper proceeds as follows. Section 2 situates the study in the live debate over AI’s incidence and in the microsimulation literature on technology shocks and household income. Section 3 describes the exposure mapping to SOC2020 major groups, the displacement, wage and capital mechanics, the four incidence families, and the PolicyEngine UK framework. Section 4 presents the distributional incidence of the central shock, the scenario grid, the five incidence families including the measured family, the constituency geography, the age and gender dimensions, benefit caseloads, robustness checks and paper-anchored scenarios. Section 5 turns to policy responses: reform counterfactuals in the shocked world and the tax-composition channel. Section 6 concludes, draws out the policy implications and sets out limitations and further work.

2 Related literature

My study sits at the intersection of two literatures: an unresolved debate about *where* in the labour market AI’s displacement effects will land, and a microsimulation tradition that traces labour-market shocks through tax-benefit systems to household disposable income. I take the incidence debate seriously as a live disagreement—each of its schools becomes a scenario family in my design—and I extend the microsimulation tradition by making incidence itself the object of variation.

2.1 Measuring exposure to artificial intelligence

The modern exposure literature descends from the occupation-level automation probabilities of Frey and Osborne (2017), whose estimate that 47 per cent of US employment was at high risk of computerisation set the template of scoring occupations against machine capabilities. Subsequent work replaced whole-occupation probabilities with task- and ability-level measurement: Webb (2019) matches patent text to job-task descriptions; Felten et al. (2021) link progress in ten AI application areas to the O*NET ability requirements of each occupation, producing the AI Occupational Exposure index I build on; and Tolan et al. (2021) connect AI benchmark results to cognitive abilities and tasks. The arrival of large language models prompted a further generation of measures targeted at generative AI specifically: the task-level rubric of Eloundou et al. (2024), the ILO’s global assessment of jobs at risk of transformation rather than elimination (ILO, 2025), the UK government’s occupational assessment (DSIT and DfE, 2023), and the GAISI index of Henseke et al. (2025). Two features of this literature matter for my design. First, whatever their construction, these indices agree that measured exposure rises with education and earnings—a reversal of the routine-biased pattern of the computerisation era (Autor et al., 2003; Goos and Manning, 2007; Goos et al., 2014). Second, exposure is not incidence: an index records what AI *could* do to an occupation’s task bundle, and translating it into who actually loses work requires assumptions the indices themselves cannot supply. The

complementarity adjustment of Pizzinelli et al. (2023), which I implement in Section 3.1, is one such assumption; the incidence families of Section 3.4 treat the remainder of that translation as an explicit scenario choice.

2.2 The incidence debate

The exposure school. The dominant approach scores occupations by the overlap between their task content and the capabilities of modern AI systems. Whether built from linkages between AI applications and occupational abilities (Felten et al., 2021), patent-to-task text matching (Webb, 2019), or expert and model assessments of which tasks large language models can perform (Eloundou et al., 2024), these indices consistently show exposure rising with education and earnings: professional, managerial and administrative occupations score highest, a pattern that recurs across European labour markets (Williamson et al., 2025) and for the UK specifically in the government’s official assessment (DSIT and DfE, 2023) and in the GAISI index of Henseke et al. (2025). On this view AI breaks with the routine-biased pattern of earlier waves, in which computerisation hollowed out codifiable mid-skill work (Autor et al., 2003; Goos and Manning, 2007; Goos et al., 2014; Acemoglu and Restrepo, 2022), and instead aims at the top of the wage distribution. The IMF strand adds an important qualification: high exposure need not mean displacement, because the same workers may enjoy the largest complementarity gains. Cazzaniga et al. (2024) formalise this ambiguity, and Pizzinelli et al. (2023) document cross-country variation in the exposure–complementarity mix; Rockall et al. (2025) apply the framework to the UK and conclude that AI could *reduce* wage inequality if displacement falls disproportionately on high earners—unless complementarity among those same workers offsets it. Two studies have carried exposure-based scenarios through national tax-benefit models—for Ireland (Doorley et al., 2026) and, for earlier automation, across European countries (Doorley et al., 2023), the latter finding that progressive taxes and means-tested transfers muted the pass-through of rising wage dispersion into household inequality. Both take exposure-proportional incidence as given; neither varies it.

The expertise-compression school. A second view, associated with Autor (2024), holds that AI’s distinctive property is not that it targets exposed occupations but that it compresses the returns to accumulated expertise. By supplying decision-relevant knowledge on demand, AI lets less experienced and less credentialed workers perform tasks previously reserved for elite professionals—an opportunity to rebuild middle-class work, but a threat to the wage premia of the experienced workers whose scarcity those premia reflect. Experimental evidence is consistent with the mechanism: generative AI raises output most for the least experienced workers, compressing within-occupation productivity gaps in customer support (Brynjolfsson et al., 2025b) and professional writing (Noy and Zhang, 2023). On this view the shock lands on senior premia rather than junior heads—nearly the mirror image of the junior-concentrated evidence below, which is precisely why incidence needs to be varied rather than assumed.

Complementarity in judgement. A third strand predicts that AI, by making prediction cheap, raises the value of the human complements to prediction—judgement, decision-making and responsibility (Agrawal et al., 2018). Workers whose roles centre on decisions rather than predictable task execution would then gain, pushing incidence away from senior decision-makers and towards workers whose tasks are prediction-like. This school shares with expertise com-

pression a scepticism that exposure indices identify the losers, while pointing the gradient in a different direction.

Adaptability and null effects. Finally, realised effects may simply be small. Danish population-scale evidence finds minimal effects of AI chatbots on earnings and hours in the first years of diffusion (Humlum and Vestergaard, 2025), and US vacancy data show exposed establishments reorganising hiring without detectable aggregate employment effects (Acemoglu et al., 2022). Acemoglu (2025) likewise derives modest macroeconomic gains from task-level foundations. These results discipline the magnitudes any simulation should entertain; my grid therefore extends down to small shocks, and my paper-anchored scenarios include effects an order of magnitude below the central calibration.

The junior-concentrated evidence. The earliest ex-post evidence on generative AI points to a gradient none of the three theoretical schools quite predicted: within exposed occupations, adjustment is concentrated on *junior* workers. Klein Teeselink (2025) link generative-AI exposure to UK firm-level records and find employment reductions of around 4.5 per cent at more exposed firms, concentrated in junior positions; Hosseini and Lichtinger (2026) document, in US resume and posting data, junior employment falling by around 9 per cent within six quarters of firm AI adoption while senior employment is unaffected; and Brynjolfsson et al. (2025a) find a 16 per cent relative employment decline among workers aged 22–25 in the most exposed US occupations since late 2022, with older workers largely spared. If incidence is graded by seniority as much as by skill, the distributional consequences differ from a classic skill-biased shock, because junior workers are spread across the household income distribution and interact differently with means-tested benefits.

2.3 From market incomes to disposable incomes

Whichever school proves right, the welfare consequences run through the tax-benefit system, and market-income incidence is a poor guide to disposable-income incidence. The European evidence on earlier automation makes the point directly: tax-benefit systems absorbed much of the routine-biased shock before it reached household inequality (Doorley et al., 2023). For AI, the factor-income channel compounds the case for full microsimulation: task-based models imply that automation lowers the labour share and channels income to capital owners (Acemoglu and Restrepo, 2018), Moll et al. (2022) show that automation raises returns to wealth itself, and Korinek and Stiglitz (2019) develop the theory of worker-replacing progress and the redistribution it calls for—with progressive capital taxation examined as a response by Bastani and Waldenström (2024). Because capital income is far more concentrated than labour income, this channel raises household inequality even where wage dispersion is unchanged, and only a model that jointly handles earnings, benefits, direct taxes and capital income can net the channels out.

2.4 Positioning

The exposure school has produced tax-benefit simulations; the other schools have produced theory and early ex-post estimates; to my knowledge no study has run the competing incidence assumptions against each other through a full tax-benefit model. That is what I do. I hold the aggregate shock fixed, implement each school’s incidence gradient as an explicit scenario family—exposure-proportional, junior-concentrated, expertise-compression, and uniform—and

let the UK tax and transfer system reveal how much the disagreement matters for poverty, inequality and the public finances.

3 Data and methods

The analysis proceeds in three stages. I first attach an occupation-level measure of exposure to artificial intelligence to each worker in a representative UK household survey. I then translate an assumed economy-wide AI shock into person-level changes to employment, wages and capital income, under alternative assumptions about who bears the displacement. Finally, I pass baseline and shocked data through an open-source tax-benefit microsimulation model to recover the distributional and fiscal consequences. This section describes each stage; the incidence scenarios that form the paper’s central axis are formalised in Section 3.4.

3.1 Measuring AI exposure

My starting point is the AI Occupational Exposure (AIOE) index of [Felten et al. \(2021\)](#), which links progress in ten AI application areas to the occupational ability requirements catalogued in O*NET, yielding a standardised score for each occupation that captures how much of its ability bundle overlaps with what AI systems can now do. Exposure alone does not distinguish occupations in which AI substitutes for workers from those in which it augments them. I therefore combine the AIOE with the complementarity adjustment of [Pizzinelli et al. \(2023\)](#), who construct an index $\theta_l \in [0, 1]$ from O*NET work-context descriptors — responsibility for others, physical proximity requirements, the consequence of error — and job-zone preparation requirements. High- θ occupations are those in which AI is likely to complement human judgement rather than displace it. The complementarity-adjusted exposure measure (C-AIOE) shields high-complementarity occupations from the raw exposure score:

$$\text{C-AIOE}_l = \text{AIOE}_l \times (1 - (\theta_l - \theta_{\min})), \quad (1)$$

where $\theta_{\min}$ is the minimum complementarity across occupational groups, so that the least complementary group retains its full exposure score. Alternative exposure measures include [Tolan et al. \(2021\)](#), the task-level rubric of [Eloundou et al. \(2024\)](#), and the cross-country applications of [Cazzaniga et al. \(2024\)](#); Section 4 reports the sensitivity of my headline results to substituting two of these for the C-AIOE.

Constructing the index for the UK requires a bespoke crosswalk, whose provenance I state in full because no published UK C-AIOE series exists at the level I need. AIOE scores are taken from the mapping to UK SOC2020 published with the DSIT analysis of AI and UK jobs ([DSIT and DfE, 2023](#)), itself derived from the [Felten et al. \(2021\)](#) scores via a chained crosswalk (US SOC2018 to US SOC2010, to ISCO-08, to UK SOC2020). The complementarity indices θ are not published at occupation level, so I reconstruct them from the public O*NET 27.3 release following the recipe documented by [Pizzinelli et al. \(2023\)](#); the reconstruction reproduces the published cross-occupation ordering up to a level offset, which is immaterial because equations (1) and (3) use θ only relative to its minimum and its mean. Both series are aggregated to the nine SOC2020 major groups using employment weights from ASHE 2025 Table 14, which covers 340 of the 412 four-digit unit groups; the remainder carry no ASHE employment estimate and are excluded from the weighting. Finally, the C-AIOE scores are shifted so that the least-exposed major group scores exactly zero. Under the allocation rule in Section 3.3 that

group receives no job losses, and the shift prevents negative values of the standardised index from corrupting the allocation weights.

The one-digit level of aggregation is a genuine constraint. The Family Resources Survey (FRS) person records I use carry only the one-digit SOC2020 code, so exposure varies across just nine major groups and all within-group variation is lost. The exposure gradient I impose on the population is correspondingly coarse, and results should be read with that in mind; imputation of finer occupation codes from the Labour Force Survey is a planned refinement.

3.2 Shock scenarios

The size of the aggregate shock cannot be estimated from historical data and must be assumed; I treat it as a scenario input and vary it systematically. The central calibration takes a displacement rate of 7 per cent of employees and an aggregate wage uplift of 2.6 per cent for those who remain in work. Both figures are calibration conventions derived from [Briggs and Kodnani \(2023\)](#): the 7 per cent converts their task-exposure estimates into an employment displacement share, and the 2.6 per cent converts their productivity projections into an aggregate wage gain. Neither is a forecast, and I do not present them as one. Alongside the labour-market effects, AI adoption is assumed to raise the return to capital by 0.4 percentage points, from a baseline of 1.005 per cent to 1.405 per cent.

Because displacement and wage growth are genuinely uncertain, the central calibration sits inside a grid spanning displacement rates from 0 to 10 per cent and aggregate wage growth from 0 to 5 per cent — 66 scenario cells in all, with the capital-return shock applied throughout. The 1 per cent “low” anchor is a pure sensitivity case: [Acemoglu \(2025\)](#) estimates a ten-year GDP effect of around 1 per cent, not an employment displacement rate, and I do not attribute the anchor to him. At the top of the grid, the 13 per cent “high” preset stress-tests a displacement rate roughly double the central convention. Section 4 additionally reports three employment-only scenarios transported from recent empirical headline estimates ([Klein Teeselink, 2025](#); [Brynjolfsson et al., 2025a](#); [Hosseini and Lichtinger, 2026](#)); these are author-designed stress tests that carry each paper’s headline magnitude into my framework, not direct implementations of those studies.

3.3 Implementing the shock at the micro level

Employment shock. The aggregate number of displaced workers is the scenario displacement rate multiplied by weighted employee employment. The universe is all employees: workers whose SOC code cannot be matched to an exposure score form a pseudo-group assigned the employment-weighted mean exposure, so no employee is mechanically excluded from risk. The total is allocated across occupational groups l in proportion to each group’s employment multiplied by its exposure:

$$\text{JobLoss}_l = \text{TotalJobLoss} \times \frac{\text{EMP}_l \times \text{C-AIOE}_l}{\sum_l \text{EMP}_l \times \text{C-AIOE}_l}, \quad (2)$$

so job losses concentrate in large, highly exposed groups and the zero-scored least-exposed group loses none. Within each group, displaced individuals are selected by ordering group members on independent uniform random keys and filling the group’s weighted quota from equation (2), with the record straddling the quota boundary included with probability equal to the remaining shortfall over its weight. Because the ordering keys are uniform, a record’s inclusion probability

does not depend on its grossing weight; the weights enter only through the quota accounting, which is the design intent.

Displacement is modelled as full withdrawal from employee work. A displaced worker’s employment income, hours, employee pension contributions and salary-sacrifice amounts are set to zero, any statutory pay is removed, and labour-market status is set to unemployed before the shocked simulation runs. Dual jobholders retain their self-employment income; only the employee job is lost. PolicyEngine UK takes no unemployment-duration input, so displaced workers are assessed as current-period unemployed under the model’s standard benefit rules rather than as long-term unemployed with contributory entitlements exhausted. To the extent that new-style (contributory) Jobseeker’s Allowance would in practice be exhausted within a year, my results may modestly overstate the benefit support reaching displaced workers; I flag this in Section 6.2.

Wage shock. Workers who survive the employment shock share an aggregate uplift pool equal to the scenario wage-growth rate times the surviving wage bill. The pool is distributed in proportion to complementarity, normalised by the employment-weighted mean complementarity over baseline workers:

$$\text{WageChange}_l = \text{WageGrowth} \times \frac{\theta_l}{\bar{\theta}} \times \text{Wage}_l, \quad \bar{\theta} = \frac{\sum_l \text{EMP}_l \theta_l}{\sum_l \text{EMP}_l}, \quad (3)$$

where EMP_l is baseline (pre-shock) employment. The normalisation ensures the aggregate uplift equals the assumed wage growth; conservation is verified in every seeded run. The economic logic mirrors equation (1): AI complements rather than substitutes for high- θ occupations, so those occupations capture a larger share of the productivity gains — an assumption the compression scenario of Section 3.4 deliberately inverts.

Capital shock. The 0.4 percentage-point increase in the return to capital is implemented by scaling each person’s savings-interest and dividend income by the ratio of shocked to baseline return, $1.405/1.005 \approx 1.398$. Rental income is excluded from the scaling, on the grounds that the AI productivity channel operates through financial rather than housing capital. The self-employed are outside both labour-market shocks; their interest and dividend income is scaled like everyone else’s.

3.4 Incidence scenarios

The rules above fix how large the shock is; they do not settle who bears it, and the empirical literature has not settled that question either. I therefore treat incidence as a first-class scenario axis. Holding the central aggregate shock fixed — 7 per cent displacement, a 2.6 per cent wage uplift, the capital shock — I re-allocate the same weighted quota of displacement under four incidence families:

Exposure-proportional. The default allocation of equations (2) and (3): displacement follows the C-AIOE gradient and wage gains follow complementarity. This family is the central calibration used throughout the shock-size grid.

Junior-concentrated. Displacement draws are tilted toward workers aged 16–24 by a multiplier equal to the ratio of junior to overall displacement effects reported by Klein Teeselink (2025) for the UK (5.8/4.5), consistent with the seniority-biased evidence of Hosseini and Lichtinger (2026) and Brynjolfsson et al. (2025a). The group quotas and wage rule are otherwise unchanged.

Expertise-compression. A stylised, author-designed stress test of the view that AI compresses expert rents rather than displacing routine work. Displacement draws are tilted (multiplier 2) toward the top weighted-earnings tertile of workers in above-median-exposure occupations, and the wage uplift is allocated by *inverse* complementarity — θ is reflected about its range, $\theta'_l = \theta_{\max} + \theta_{\min} - \theta_l$, so that mid-skill workers capture the capability gains — with the same employment-weighted normalisation as equation (3). I emphasise that this is a stress test, not a calibrated implementation of any published estimate.

Uniform. Every employee is assigned equal exposure and equal complementarity, so displacement risk and wage gains are flat across occupations. This family isolates the contribution of the incidence gradient itself: differences between it and the others are attributable entirely to who bears the shock.

Each family delivers the same weighted count of displaced workers to within rounding, so cross-family differences in fiscal cost, poverty and inequality measure the consequences of incidence alone. The incidence gradient in each family is an assumption, not an estimate; the paper’s contribution is to show how the tax-benefit system mediates whichever incidence materialises.

Measured-incidence family. Section 4.5 adds a fifth family that replaces the stylised gradients above with the displacement patterns measured in UK firm-level data by Klein Teeselink (2025). Displacement probability is proportional to the shifted C-AIOE multiplied by three tilts: a junior multiplier of 1.29 (the ratio of their junior to overall displacement effects, 5.8/4.5), a London multiplier of 1.5, and a wage-tier multiplier of 9.6 above versus 0.6 below weighted-median earnings, reflecting their finding that losses concentrate in high-paying firms; the aggregate shock stays at the central preset. Provenance discipline matters here: the junior ratio and the direction of the London and pay-tier tilts come from Klein Teeselink (2025), but the London magnitude of 1.5 (their finding is qualitative), the low-wage floor of 0.6, the proxying of firm pay tiers by person earnings, the use of age under 25 as a proxy for their seniority concept, and the multiplicative independence of the three channels are all author-imposed. The underlying parameters were moreover taken from the accompanying press release and secondary summaries, the SSRN full text being inaccessible at the time of writing; each run’s output records this provenance split explicitly.

Policy reforms in the shocked world. The reforms of Section 5 are evaluated as $M(\text{shock} + \text{reform}) - M(\text{shock})$ on the fixed seed-0 shock draw, applying for 2026 only, with poverty lines held at their shocked-world values. Benefit-side reforms are implemented as PolicyEngine UK parameter reforms — scaling the Universal Credit standard allowance, raising the benefit-cap thresholds and cutting the taper rate — so their costs are net of the means-testing clawback the model computes. Wage insurance is instead a direct post-simulation transfer to displaced workers (half of lost earnings, capped), treated as non-taxable and disregarded by means tests, so its gross cost equals its net cost and brackets the generous end of possible implementations.

Constituency geography. The constituency results of Section 4.6 rest on PolicyEngine’s $650 \times 53,508$ weight matrix, which assigns each household in the *enhanced* FRS 2023–24 dataset a calibrated weight in every Westminster parliamentary constituency. The SERNUM-based occupation join of Section 3.5 is invalid on that dataset, so SOC major group is imputed from the plain-FRS 2024–25 weighted distribution within age band $\times$ gender $\times$ region $\times$ earnings-decile cells (99.4 per cent weighted coverage from the full four-way cell, seed 0). Constituency variation therefore reflects demographic and earnings composition plus a single displacement draw, not observed occupation mix, and the geographic results — computed on a different dataset and period — are not directly comparable to the poverty numbers elsewhere in the paper.

3.5 Microsimulation with PolicyEngine UK

I use PolicyEngine UK (version 2.89.2), an open-source tax-benefit microsimulation model, running on person-level microdata from the Family Resources Survey 2024–25 obtained from the UK Data Service. Occupation codes are not carried in the processed PolicyEngine dataset, so I join the SOC2020 major group from the raw FRS `adult.tab` file onto the simulation’s person table using the identifier convention `person_id = SERNUM  $\times$  1000 + PERSON`, which matches every FRS `adult`. The simulation year is 2026.

The counterfactual design is a paired comparison on identical data: a baseline simulation and a shocked simulation are run on the same records, with the shocked run receiving the modified employment income, hours, pension contributions, savings-interest income, dividend income and employment status of Section 3.3 as direct inputs. All reported effects are differences between the two runs. This design removes between-run noise — the survey sample and grossing weights are identical on both sides, so no part of a reported effect reflects resampling — but it does not quantify population sampling uncertainty: both runs inherit whatever sampling error the FRS itself carries, and my Monte Carlo intervals speak only to the stochastic identity of the displaced, not to survey sampling.

All income results are expressed on the HBAI cash disposable income concept — equivalised household net income as defined for the Households Below Average Income statistics and implemented in PolicyEngine UK — which is also the concept underlying the poverty lines. I report four outcome families: the change in the government balance (the Exchequer cost of the shock, net of additional revenue on higher wages and capital income); person-weighted poverty rates before and after housing costs (BHC and AHC); the Gini coefficient of equivalised household disposable income, computed with negative incomes bottom-coded at zero; and mean changes in household income by decile, with deciles fixed at their baseline positions so that individuals are tracked rather than re-ranked.

Because the identity of the displaced is stochastic, I adopt explicit seed conventions and label every result accordingly. The 66-cell grid, the three presets, the decomposition figure and the four incidence families are single draws at seed 0, so that cells and families differ only in their scenario inputs, never in their random draw. The decile transition and wage profiles (Figures 1 and 2) average 50 seeded draws. The robustness summary and the decomposition confidence band use a 20-draw Monte Carlo at seeds 0–19; wherever a $\pm$ value appears it is the standard deviation across draws, not a survey-based standard error.

4 Results

I present results in eight steps. Sections 4.1 and 4.2 trace the central scenario from the three market-income channels through the tax-benefit system to household disposable income. Section 4.3 reports the fiscal and distributional aggregates for the preset scenarios, with Monte Carlo variation. Section 4.4 widens the lens to a 66-cell grid over shock sizes. Section 4.5 then turns to the paper’s central question—who bears the shock—by holding the aggregate shock fixed and varying its incidence across five families, four stylised and one shaped by measured UK evidence. Section 4.6 maps the shock across parliamentary constituencies, Section 4.7 examines the age and gender dimensions, and Section 4.8 translates the shock into benefit caseloads. Throughout, all income and inequality results are computed on the HBAI cash disposable income concept described in Section 3.5. Seed and draw conventions follow Section 3.5 and Appendix A.1: the grid, the decomposition, the preset table and the incidence families are single draws at seed 0, the decile profiles average 50 draws, and $\pm$ values are standard deviations across 20 Monte Carlo draws.

4.1 The three market-income channels

The central scenario displaces 7 per cent of exposed employment—1.61 million workers—and grants surviving workers in exposed occupations a wage uplift averaging 2.6 per cent, alongside the capital income channel. Figure 1 plots the share of each equivalised disposable income decile that transitions from employment to unemployment. The gradient is strictly monotone: 0.50 per cent of the bottom decile transitions, rising through every decile to 5.05 per cent at the top. It is important to be clear about what this figure shows. The gradient is not an empirical finding about who AI will in fact displace; it is the exposure-proportional allocation assumption made visible. Because the professional, associate-professional and administrative occupations that score highest on the C-AIOE index sit in the upper half of the income distribution, allocating displacement in proportion to occupational exposure necessarily concentrates job loss among higher-income households. The independent Irish estimates of Doorley et al. (2026), built on different microdata and a finer occupational classification, display a similarly top-concentrated gradient, which suggests the pattern is a general property of services-intensive economies under exposure-proportional incidence—but the incidence itself remains an assumption, and Section 4.5 treats it as one.

The wage channel, by contrast, is almost distributionally flat. Figure 2 shows the percentage change in employment income among workers who remain employed: gains run from 2.51 per cent in the bottom decile to 2.72 per cent in the top, a spread of barely 0.2 percentage points around the 2.6 per cent average uplift. Exposure is sufficiently widespread across the deciles in which people work that the productivity dividend is spread thinly and evenly. It is the displacement margin, not the wage margin, that drives the distributional results below.

The third channel is capital income. The scenario assumes part of the productivity gain is capitalised into returns on existing capital, applying a uniform proportional uplift to interest and dividend income (Section 3.3). Although the percentage change is identical across deciles by construction, Figure 3 shows the incidence in pounds is heavily concentrated: the top decile holds 35.2 per cent of all capital income (£8.1 billion of the baseline aggregate in my data), the top two deciles together more than half, and the bottom decile just 2.8 per cent. Whatever share of AI gains is capitalised flows overwhelmingly to households already at the top of the distribution.




Figure 1: Share of the population transitioning from employment to unemployment by equivalised disposable income decile, central scenario. Bars are means over 50 seeded displacement draws; error bars show 95 per cent confidence intervals across draws. The gradient is the exposure-proportional allocation assumption made visible, not an empirical incidence estimate.




Figure 2: Percentage change in employment income among workers remaining in employment, by income decile, central scenario. Means over 50 seeded draws.




Figure 3: Capital income channel by decile: a uniform proportional uplift, but pound gains concentrated where capital income is held.

4.2 Household mediation of the shock

Figure 4 decomposes the change in household income by decile into market income, benefits, and taxes and contributions, all expressed relative to baseline disposable income. The decomposition is household-level and additive: market income change plus benefit change minus tax change equals the disposable income change up to an explicit residual, which collects perimeter items inside the HBAI concept (council tax and pension deductions) and never exceeds 0.5 per cent of baseline income in any decile.

Market income losses follow the transition gradient: they are near zero in the bottom decile, around 1–1.6 per cent of baseline disposable income in deciles 3–5, and reach 6.2 and 5.3 per cent in deciles 9 and 10. The tax-benefit system then cushions these losses through two distinct mechanisms. Across the middle deciles, automatic benefit entitlements—Universal Credit in particular—rise by up to 0.5 per cent of baseline income as displaced earners qualify for means-tested support. At the top the cushioning operates through the tax system instead: deciles 9 and 10 see their income tax and National Insurance liabilities fall by 1.3–1.4 per cent of baseline income as high-marginal-rate earnings disappear. The combined offset is substantial. In decile 9, a market income loss of 6.2 per cent of baseline disposable income becomes a disposable income loss of 4.3 per cent—roughly three-tenths of the market shock absorbed—and in decile 5 a 1.6 per cent market loss becomes a 1.0 per cent disposable loss. Averaged over 20 displacement draws, disposable income losses run from 0.53 per cent in the bottom decile to 4.09 per cent in the top; the series is not perfectly monotone—decile 2 loses slightly more than decile 3 (−0.88 against −0.82 per cent)—but the broad upward gradient is stable across draws.

4.3 Fiscal and distributional aggregates

Table 1 summarises the preset scenarios. In the central scenario the annual Exchequer cost is £18.4 billion, driven by lost income tax and National Insurance on displaced higher earners to-




Figure 4: Additive household-level decomposition of the change in income by decile into market income, benefits, and taxes and contributions, central scenario (single draw, seed 0), as a share of baseline HBAI disposable income. The residual collects perimeter items (council tax, pension deductions) inside the HBAI concept.

gether with higher benefit spending. Poverty rises by 1.87 percentage points before housing costs (1.88 points after housing costs) and the Gini coefficient of equivalised disposable income rises by 1.10 percentage points from a baseline of 0.309. The low scenario—a 1 per cent displacement rate with no wage gain, which I treat purely as a sensitivity case rather than an estimate drawn from the literature—costs £2.5 billion and moves poverty and the Gini by less than 0.2 points each. The high scenario, at 13 per cent displacement, costs £46.7 billion a year with poverty up 3.80 points. PolicyEngine’s before-housing-costs poverty line is an *absolute* threshold, so these poverty changes measure absolute impoverishment rather than movement relative to a shock-shifted median.

The scenario anchors deserve a note of caution. The central calibration converts the task-exposure and productivity figures of Briggs and Kodnani (2023) into a displacement rate and a wage uplift via the conventions described in Section 3.2; it is a calibration convention, not a forecast. The low scenario is likewise not an employment estimate from Acemoglu (2025), whose 0.9–1.1 per cent figure refers to ten-year GDP growth, not job displacement.

Table 1: Summary of preset scenarios (single draw, seed 0)

Scenario	Displaced (million)	Exchequer (£bn/yr)	Δ Poverty (BHC, pp)	Δ Gini (pp)
Low	0.23	2.5	+0.14	+0.19
Central	1.61	18.4	+1.87	+1.10
High	3.00	46.7	+3.80	+1.83
Central, junior-concentrated	1.61	19.4	+1.86	+0.92

Note: Low: 1 per cent displacement, no wage gain (sensitivity case). Central: 7 per cent displacement, 2.6 per cent wage gain. High: 13 per cent displacement, 2.6 per cent wage gain. Junior-concentrated: central shock with displacement draws tilted toward ages 16–24. The capital-return shock (+0.4pp) applies in every preset.

None of this depends on the particular displacement draw. Repeating the central scenario over 20 independent draws gives a mean Exchequer cost of £18.0 billion (standard deviation £1.5 billion, range £15.8–21.1 billion), a poverty increase of $+1.80 \pm 0.11$ percentage points and a Gini increase of $+1.04 \pm 0.09$ points. Nor does it depend on the exposure index: rerunning the central scenario with the DSIT AIOE mapping or the GPT-exposure measure of [Eloundou et al. \(2024\)](#) in place of the C-AIOE moves the Exchequer cost by less than £0.2 billion and leaves the transition gradient essentially unchanged (top-decile transition shares of 4.9 and 4.7 per cent against 4.7 under the C-AIOE; bottom-decile shares of 0.4–0.5 per cent under all three).

4.4 The shock-size grid

To move beyond the presets I simulate a grid of 66 scenarios, crossing eleven displacement rates (0 to 10 per cent of exposed employment) with six wage uplifts (0 to 5 per cent), each combined with the capital income channel, all under exposure-proportional incidence. Figure 5 maps the change in average household disposable income. The corners bracket the range: with a 5 per cent wage gain and no job loss, average disposable income rises by 3.0 per cent; with 10 per cent displacement and no wage gain it falls by 5.2 per cent. The near-zero contour runs close to the grid diagonal—roughly one percentage point of wage growth offsets one percentage point of displacement in aggregate disposable income terms—so whether AI raises or lowers average living standards depends almost entirely on the balance struck between these two forces, about which the empirical literature remains genuinely uncertain.



Figure 5: Change in average household disposable income (per cent) across the 66-cell grid of displacement rates and wage uplifts (single draw, seed 0).

Figure 6 shows the corresponding fiscal picture. One measurement issue deserves note: in PolicyEngine’s UK aggregates, income tax plus National Insurance minus total benefit expenditure (which includes the state pension) is a negative number, so expressing revenue changes relative to that net position produces uninformative percentages. I therefore express the net revenue change relative to gross income tax and National Insurance receipts. On that basis the grid spans a gain of £22.3 billion, or 7.2 per cent of receipts (5 per cent wage growth, no displacement), to a loss of £31.2 billion, or 10.0 per cent of receipts, in the worst cell (10 per

cent displacement, no wage growth). At the central displacement rate, 7 per cent with no wage offset costs £21.5 billion (6.9 per cent of receipts), which a 3 per cent wage uplift roughly halves to £10.2 billion. As with disposable income, the break-even locus lies close to the diagonal.




Figure 6: Net Exchequer revenue change as a percentage of gross income tax and National Insurance receipts, across the 66-cell scenario grid.

Figure 7 maps the change in the Gini coefficient across the same grid, and delivers the paper’s most robust result: the Gini rises in all 66 cells, from +0.06 percentage points in the capital-only cell (no displacement, no wage gain) to +1.55 points at the far corner. No combination of parameters within the grid reduces measured inequality. The mechanism is one of polarisation rather than a simple rich-get-richer story, and it coexists with an incidence that is itself progressive: displacement strikes hardest at the top of the earnings ladder and pushes affected households down the distribution, while the wage uplift lifts the surviving employed—disproportionately in the upper-middle and top deciles—further away from workless and low-income households, and the capital channel reinforces the very top. Because both the loss side and the gain side of the shock widen gaps between households, inequality rises even though the market shock lands mainly on higher earners. That tension—a top-loaded shock that still raises the Gini—recurs throughout what follows.

4.5 The incidence axis: who bears the shock

Everything so far conditions on exposure-proportional incidence. But the allocation of a given aggregate shock across workers is at least as uncertain as its size, and the empirical literature points in different directions: exposure indices imply top-loaded incidence, the hiring-margin evidence points to juniors (Brynjolfsson et al., 2025a; Hosseini and Lichtinger, 2026), and the expertise-compression view suggests AI erodes the premium of mid-skill judgement work. I therefore treat incidence as a first-class scenario axis. Holding the aggregate shock fixed at the central calibration—the same 1.6 million displaced workers, the same wage and capital channels—I reallocate displacement across five families. Four are stylised: *exposure-proportional* (the central calibration used above); *junior-concentrated*, which tilts displacement towards early-career workers; *expertise-compression*, an author-designed stress test in the spirit of the task literature



Figure 7: Change in the Gini coefficient of equivalised HBAI disposable income (percentage points) across the 66-cell scenario grid. The Gini rises in every cell.

descending from [Autor et al. \(2003\)](#), which concentrates displacement on mid-to-upper-skill occupations whose comparative advantage lies in codified expertise; and *uniform*, which allocates displacement occupation-neutrally. The fifth, *measured*, shapes displacement probabilities by the UK evidence of [Klein Teeselink \(2025\)](#): exposure-proportional draws multiplied by a junior tilt (1.29), a London tilt (1.5) and a wage-tier tilt that loads displacement onto workers above median earnings (9.6 above, 0.6 below), while the aggregate shock stays at the central preset. Table 2 and Figure 8 report the results (single draw, seed 0).

Table 2: The same aggregate shock under five incidence families (central calibration, single draw, seed 0)

Incidence family	Displaced (million)	Exchequer cost (£bn/yr)	Δ Poverty BHC (pp)	Δ Poverty AHC (pp)	Δ Gini (pp)
Exposure-proportional	1.61	18.4	+1.87	+1.88	+1.10
Junior-concentrated	1.61	19.4	+1.86	+1.83	+0.92
Expertise-compression	1.61	24.5	+2.12	+2.08	+0.92
Uniform	1.62	14.2	+1.83	+1.86	+1.14
Measured (Klein Teeselink)	1.61	32.5	+2.29	+2.61	+1.02

Note: The measured family combines the exposure-proportional draw with a junior multiplier (1.29), a London multiplier (1.5) and wage-tier multipliers (9.6 above / 0.6 below median earnings). Its quantitative parameters were taken from the KCL press release and secondary summaries of [Klein Teeselink \(2025\)](#), pending verification against the SSRN full text; the London multiplier, the low-wage floor, the multiplicative combination of channels and the age-under-25 proxy for junior positions are author-imposed.

Three findings stand out. First, who bears the shock determines what it costs. The Exchequer cost of an identical aggregate displacement ranges from £14.2 billion under uniform incidence to £32.5 billion under the measured family—a spread of £18 billion a year, roughly the central estimate itself, generated purely by reallocating which workers lose their jobs. Top-loaded incidence is fiscally expensive because displaced high earners take large income tax and




Figure 8: Decile incidence of the same aggregate central shock under five incidence families: share of each decile transitioning to unemployment, and change in equivalised HBAI disposable income by decile (single draw, seed 0).

National Insurance payments with them; uniform incidence is cheaper because uniformly drawn displaced workers earn, and therefore pay, less.

Second, the fiscal saving from flatter incidence buys almost no poverty relief. Poverty rises by 1.83 to 1.87 points before housing costs in three of the five families; the uniform shock, despite costing the Exchequer £4.2 billion less than the exposure-proportional one, leaves poverty essentially as high (+1.83 against +1.87 points). Cheap incidence is not benign incidence: the fiscal and poverty margins move separately, and a government cannot infer from a modest revenue hit that households near the poverty line have been spared.

Third, the top-loaded families are the worst of both worlds. Expertise compression, a stylised stress test of my own design rather than an implementation of any published estimate, costs £24.5 billion and raises BHC poverty by +2.12 points. The measured family goes further on both margins: at £32.5 billion it is the costliest of the five, and its poverty increase (+2.29 points BHC, +2.61 points AHC) is the largest. Its wage-tier multiplier concentrates displacement on workers above median earnings—the most top-tilted decile gradient of the five families, with 0.02 per cent of the bottom decile transitioning against 6.6 per cent of the top—so it maximises the loss of income tax and National Insurance, while the affected households’ fall from the middle and upper-middle of the distribution drags more households across the (absolute) poverty line than a purely top-decile shock would. To the extent that the measured UK evidence is a better guide to incidence than occupational exposure alone, my central preset understates both the fiscal and the poverty cost of the shock.

The Gini, meanwhile, rises in all five families—by +0.92 to +1.14 percentage points—just as it rises in all 66 grid cells. Incidence reallocates the cost between the fiscal and poverty margins; it does not change the direction of the inequality effect.

4.6 The geography of the shock

The incidence question has a spatial counterpart: an occupationally selective shock lands unevenly across places because occupations cluster geographically. To map it I project the central

scenario onto the 650 Westminster parliamentary constituencies using PolicyEngine’s calibrated constituency weight matrix. One methodological caveat governs everything in this subsection. The weight matrix indexes the *enhanced* FRS 2023–24 dataset, on which occupation codes cannot be joined directly, so SOC major group is imputed from the plain-FRS distribution within age-band $\times$ gender $\times$ region $\times$ earnings-decile cells (99.4 per cent weighted coverage from the full four-way cell). Constituency variation therefore reflects demographic and earnings composition under a single seed-0 draw, not observed local occupation mixes, and the results are not directly comparable to the poverty numbers elsewhere in the paper: on this dataset and period the national BHC poverty response is +3.9 percentage points, roughly twice the +1.87 points of the central run, a difference attributable to the dataset (enhanced 2023–24 versus plain 2024–25) and simulation period.

With that understood, the geography is far from uniform. Nationally, aggregate household income falls by 1.03 per cent and 71.5 workers per 1,000 are displaced. Across the twelve regions the income hit is deepest in Northern Ireland (–2.29 per cent) and Scotland (–1.79 per cent) and mildest in Wales (–0.78 per cent) and the South East (–0.83 per cent), even though displacement rates span only 61 (Northern Ireland) to 79 (Wales) per 1,000 workers: where displacement is similar, the income consequences depend on which earners are hit and how much cushioning the local benefit mix provides. At constituency level the spread widens further. The worst-hit constituencies are Lagan Valley (–5.43 per cent), Upper Bann (–4.11 per cent) and Belfast East (–4.03 per cent), with Stratford and Bow and Islington South and Finsbury close behind—a pattern combining Northern Irish seats with high-exposure inner-London ones—while a small number of constituencies gain on net, led by Kingston and Surbiton (+2.05 per cent) and Dumfriesshire, Clydesdale and Tweeddale (+1.75 per cent), where wage and capital gains outweigh displacement losses. Figure 9 maps both surfaces.

4.7 The age and gender dimensions

The incidence axis has a within-family age structure worth drawing out. Under exposure-proportional incidence, workers aged 16–24 account for only 4.9 per cent of the displaced in the central scenario, because young people in the UK are concentrated in low-exposure service occupations; the burden falls on prime-age workers (28.1 per cent of the displaced are aged 35–44). The junior-concentrated family of Table 2 reweights displacement probabilities towards junior workers while holding the aggregate displacement rate at 7 per cent, raising the 16–24 share to 6.2 per cent and lowering the 55–64 share from 19.1 to 16.2 per cent. The aggregate consequences are modest—the Exchequer cost rises slightly to £19.4 billion, poverty is essentially unchanged (+1.86 against +1.87 points) and the Gini rise falls to +0.92 points—but the exercise makes a structural point: seniority bias must be imposed, because it cannot be derived from occupational exposure alone. Even under the junior tilt, prime-age workers remain the largest displaced group.

Table 3 pushes further, reporting three employment-only stress tests anchored to headline estimates in the recent empirical literature: the UK occupational estimates of Klein Teeselink (2025), the seniority-biased hiring evidence of Hosseini and Lichtinger (2026), and the early-career employment declines in AI-exposed occupations documented by Brynjolfsson et al. (2025a). These are author-designed transports of headline numbers onto the UK employment structure, not direct implementations of the source papers’ designs, and the footnotes flag the strongest assumptions.

The lesson is that cohort-specific shocks are fiscally small. The Hosseini and Lichtinger




Figure 9: Constituency hex maps of the central scenario: change in aggregate household income (left) and displaced workers per 1,000 (right), enhanced FRS 2023–24 basis with imputed occupations, single draw (seed 0). Not directly comparable to the main-paper poverty numbers.

Table 3: Paper-anchored displacement stress tests, UK translation (employment channel only, single draw, seed 0)

Scenario	Displaced (million)	Exchequer (£bn/yr)	Δ Poverty (BHC, pp)	Δ Gini (pp)
Klein Teeselink (2025) ^a	1.03	22.5	+1.30	+0.42
Hosseini–Lichtinger (2026) ^a	0.10	1.3	+0.11	+0.04
adopter-share scaled	0.02	0.1	+0.01	+0.01
Brynjolfsson et al. (2025a) ^{a,b}	0.02	0.2	+0.02	+0.01

Note: Klein Teeselink: −4.5 per cent of employees, junior-tilted draws. Hosseini–Lichtinger: 9 per cent of under-30s in above-median-exposure occupations (at-adopter rate); the scaled row multiplies by the adopter employment share. Brynjolfsson et al.: 16 per cent of ages 22–25 in top-quintile-exposure occupations. ^aSource estimates are firm-conditional or relative declines; transporting them economy-wide (or as absolute rates) makes these upper bounds. ^bA *relative* decline versus less-exposed groups in the source paper.

(2026) and Brynjolfsson et al. (2025a) stress tests, which confine displacement to entry-level and early-career workers, displace 0.10 million and 0.02 million workers and cost the Exchequer £1.3 billion and £0.2 billion a year respectively—one to two orders of magnitude below the broad-based scenarios—precisely because junior workers contribute little income tax and National Insurance. Only the economy-wide Klein Teeselink (2025) transport, with just over one million displaced, generates fiscal effects comparable to my central preset (£22.5 billion, poverty +1.30 points). If AI adoption in practice falls hardest on juniors through recruitment rather than dismissal, that operates through hiring channels that stock-based simulation frameworks such as mine do not capture by construction.

Exposure also carries a gender gradient. Employed women in the FRS have a mean C-AIOE of 0.33 against 0.18 for employed men, reflecting women’s concentration in administrative, professional and associate-professional work. In the central scenario this translates into a displacement rate of 7.5 per cent for employed women against 6.4 per cent for men (means over 20 draws): women hold 51.8 per cent of employment but account for 55.6 per cent of the displaced. Household-level income changes are nevertheless similar by sex (−3.0 per cent for women, −2.8 per cent for men). A natural reading is that income pooling within couples mutes the individual asymmetry at the household level, though I have not decomposed the result and offer this as a hypothesis rather than a finding. The individual-level asymmetry itself echoes the international evidence that female employment is substantially more exposed to generative AI than male employment (ILO, 2025), and suggests that within-household earnings shares—and with them, financial autonomy—would shift even where household totals move little.

4.8 Benefit caseloads

The fiscal aggregates above can also be read on the delivery side, as caseloads. In the central scenario the shock draws 434,000 benefit units into Universal Credit while 58,000 exit, a net caseload increase of 376,000 benefit units, and annual UC spending rises by £3.3 billion from a baseline of £41.4 billion, of which £1.1 billion is the housing element within paid awards. Against DWP’s roughly £334 billion of welfare spending in 2025–26 (a context figure, not a model output), the central increment is about 1 per cent; the high scenario roughly doubles it (695,000 net new benefit units, +£6.6 billion), while the low scenario adds only 58,000 benefit units and £0.4 billion. Caseload responses vary with incidence in the expected direction: the junior-concentrated family, whose displaced workers more often live in households with other earners, adds fewer benefit units (302,000 net) at lower cost (+£2.5 billion). Two accounting notes apply: only the benefit-unit figures are net of exits—the corresponding household and person counts (415,000 and 1.0 million in the central scenario) are gross new entrants—and the housing element is measured within paid awards, slightly overstating housing-specific cash.

5 Policy responses

If the shock modelled above is structural rather than cyclical—a permanent reallocation of labour demand rather than a downturn that unwinds—then the question facing the state is not whether the existing tax-benefit system cushions the blow (Section 4 shows that it does, absorbing roughly half of the market-income loss) but whether it cushions it well, and at what fiscal cost the residual hardship could be bought down. This section evaluates three stylised policy responses inside the shocked world. Each reform is simulated on top of the central scenario (7 per cent displacement, seed 0, 2026), so the estimand throughout is the difference

between the shocked-plus-reform and shocked simulations: what the reform adds, given that the shock has happened. I make no claim about the merits of these instruments in the absence of the shock, and the reforms apply for 2026 only—they are circuit breakers, not permanent settlements.

5.1 Three stylised reforms

The first reform (R1) is a wage-insurance scheme in the spirit of trade-adjustment proposals: displaced workers receive 50 per cent of their lost earnings, capped at £15,000 per year. I implement it as a post-simulation transfer that is non-taxable and disregarded by means tests, so its gross cost equals its net cost; this brackets the generous end of plausible implementations. The second (R2) is a demand-side circuit breaker within the existing architecture: a 20 per cent increase in the Universal Credit standard allowance, implemented as a parameter reform inside PolicyEngine so that all interactions with tapers, caps and passported entitlements are captured. The third (R3) suspends the benefit cap and cuts the UC earnings taper from 55 to 45 per cent, targeting the households for whom the existing system’s own restraints bind hardest. R2 and R3 costs are net of clawback through the rest of the system.

Table 4: Policy reforms on top of the central shock, 2026 (exposure-proportional incidence, seed 0)

Reform	Cost (£bn)	Δ poverty (pp)		Δ Gini (pp)	£bn per pp averted	
		BHC	AHC		BHC	AHC
R1 Wage insurance	21.2	−1.53	−1.65	−0.97	13.8	12.8
R2 UC allowance +20%	5.8	−0.64	−0.56	−0.31	9.1	10.4
R3 Cap suspension + taper cut	4.7	−0.46	−0.45	−0.23	10.2	10.3

Note: All effects are shocked-plus-reform minus shocked, fixed seed-0 draw, reforms in force for 2026 only. R1: 50 per cent of lost earnings, £15,000 cap, paid to 1.61 million displaced workers (mean transfer £13,173), modelled as non-taxable and means-test-disregarded so gross cost equals net cost. R2 and R3 are PolicyEngine parameter reforms; costs are net of clawback. R3 combines benefit-cap suspension with a UC taper cut from 55 to 45 per cent. Poverty lines are held at their shocked-world values.

Table 4 reports the results and Figure 10 plots the decile pattern of the gains. The headline comparison is one of cost-effectiveness. The UC uplift averts a percentage point of after-housing-costs poverty for £10.4 billion, and the cap-and-taper package for £10.3 billion, against £12.8 billion for wage insurance. The reason is mechanical but instructive: wage insurance is earnings-proportional, and because the exposure-proportional shock displaces workers who are disproportionately in the upper half of the distribution, roughly 51 per cent of the R1 benefit is experienced in baseline deciles 8–10 (measured as a person-weighted household-gain share rather than a strict pound share). The UC-side instruments, by contrast, concentrate 71 per cent (R2) and 67 per cent (R3) of their benefit in deciles 1–5. Wage insurance does buy roughly three times more total poverty reduction than either UC reform—it is by far the largest intervention—but at roughly four times the cost, and with the weakest targeting per pound.

Because Section 4.5 established that incidence is the paper’s central uncertainty, I re-run the wage-insurance reform under the junior-concentrated and uniform incidence families. The cost-effectiveness ranking is robust: R1 costs £14.3 billion per AHC point under junior incidence and £11.8 billion under uniform incidence, so even in its most favourable incidence world wage insurance only just reaches the £10.3–10.4 billion achieved by the UC-side instruments under




Figure 10: The three policy reforms on top of the central shock: fiscal cost, poverty and Gini effects, and the decile distribution of household gains. All panels show shocked-plus-reform minus shocked, seed 0, 2026.

the central family. The gross cost of R1 is nearly invariant across families (£20.3–21.4 billion), as it must be when the aggregate displaced wage bill is held approximately fixed; what varies is how much of that spending lands near the poverty line.

None of this settles the choice of instrument. Wage insurance addresses a different objective—consumption smoothing and earnings-loss compensation for the displaced themselves, most of whom are not poor—and a government persuaded by that objective would not be deterred by an unfavourable poverty-targeting ratio. What the exercise shows is narrower: if the criterion is poverty averted per pound in the shocked world, the existing means-tested architecture, temporarily made more generous, dominates a new earnings-linked scheme, precisely because the shock modelled here falls disproportionately on households that means tests are designed to exclude.

5.2 The tax-composition channel

The reforms above spend money into a fiscal position that the shock has already damaged, and the size of that damage depends on a question the microsimulation cannot answer by itself: where does the displaced wage bill go? The OBR’s recent analysis of AI-related fiscal risks (OBR, 2026) emphasises that labour displacement erodes the income-tax and National Insurance base even if aggregate output is unchanged, with the fiscal outcome hinging on how much of the lost labour income reappears as taxable profit. I price this channel on top of the microsimulation with a simple accounting exercise. Let ϕ denote the share of the displaced wage bill that reappears as corporate profits taxed at the 25 per cent main rate. In the central scenario the displaced wage bill is £77.9 billion and the effective income-tax-plus-NICs rate on displaced workers’ earnings is 25.0 per cent, so the labour tax forgone is £19.5 billion—and because that effective labour rate almost exactly equals the corporation-tax main rate, the shortfall is very nearly $(1 - \phi) \times 0.25 \times$ the displaced wage bill. At $\phi = 0$ the net revenue shortfall is the microsimulation’s £21.5 billion; at $\phi = 0.5$ it falls to £11.8 billion; and at $\phi = 1$ the labour-to-capital channel is almost exactly self-financing (a £0.03 billion surplus relative to full labour taxation), leaving only a residual

£2.1 billion net cost, chiefly shock-induced benefit spending together with smaller non-labour-tax revenue changes. Figure 11 traces the full grid over displacement rates and ϕ .




Figure 11: Net revenue change by displacement rate and ϕ , the share of the displaced wage bill reappearing as corporate profits taxed at 25 per cent. The $\phi = 0$ line is the microsimulation’s net revenue change; higher ϕ adds recouped corporation tax outside the simulation.

Revenue neutrality, however, is not distributional neutrality. To see who receives the re-composed income, I simulate a dividend-recycling case at $\phi = 0.5$ with a payout ratio of 0.5 (a parameter, not an estimate): £19.5 billion of post-CT profits are distributed pro rata to existing FRS dividend holders, with the dividend tax on those distributions captured inside the simulation. The result is stark. The Gini rises by a further 0.62 percentage points relative to the shocked world without recycling, and poverty—before or after housing costs—moves by exactly zero: the recycled dividends reach no household near the poverty line. The decile pattern of the gains is lumpy because FRS dividend records are sparse, but the message is not: deciles 1–4 receive exactly nothing. Even in the most fiscally benign reading of the shock, in which the Exchequer recovers essentially all of the lost labour tax through the corporate channel, the compositional shift from wages to capital income leaves the bottom of the distribution no better off and the overall distribution more unequal—which is precisely the configuration in which the transfer-side instruments of Section 5.1 would be needed most.

Two caveats bound this exercise. It is static accounting layered on top of the microsimulation: the incidence of corporation tax is not modelled, the corporation-tax base is assumed equal to accounting profit (no allowances, losses or profit shifting), and the effective labour rate is a pro-rata attribution of income tax and NICs to employment income that ignores the schedule ordering of savings and dividend bands. And ϕ itself is the free parameter the whole channel turns on; I present a grid rather than a point estimate because nothing in my framework pins it down.

6 Conclusions and policy implications

The results should be read as scenario analysis, not prediction. I do not forecast the pace or extent of AI adoption in the United Kingdom; I trace the distributional and fiscal consequences that would follow *if* labour-market shocks of a given size and occupational incidence were to materialise, under tax-benefit rules as legislated. The early observational record justifies the full width of the grid, including its benign corner: population-scale Danish evidence finds minimal earnings and hours effects of chatbot adoption to date (Humlum and Vestergaard, 2025), and postings-based analyses disagree on whether the entry-level slowdown in exposed occupations is attributable to AI at all. UK postings evidence points to a modest relative contraction in exposed occupations since late 2022 (Henseke et al., 2025), but nothing yet at the scale of my central scenario. The Office for Budget Responsibility treats AI principally as a productivity risk to its fiscal projections (OBR, 2025); my estimates supply the complementary labour-incidence arithmetic that such projections currently leave implicit. The scenario parameters are calibration conventions transported from headline estimates in the empirical and modelling literature—including Briggs and Kodnani (2023) and Cazzaniga et al. (2024)—rather than elasticities estimated on UK data, and the wide span between my mildest and harshest cells is a deliberate acknowledgement of the uncertainty surrounding all of these anchors. The incidence gradient itself is an assumption in every family I run; what the paper establishes is how the tax-benefit system mediates each assumed incidence.

Two findings organise the discussion. The first is that the UK tax-benefit system provides substantial automatic insurance to low-income households: means-tested support—principally Universal Credit—absorbs a large share of the market-income losses in the lower deciles, while progressive income taxation claws back part of the losses higher up before they reach disposable income. This cushioning is not free. In the central scenario the combined cost of lower receipts and higher benefit expenditure is £18.4 billion a year, and the cost scales with the severity of the shock. The protection households experience is purchased through a deterioration of the public finances.

The second finding is that the Gini of equivalised household disposable income rises in every one of the 66 shock-size cells and in all five incidence families (Sections 4.4 and 4.5), even though the market-income incidence in the exposure-proportional family is progressive—displacement risk rises monotonically with household income (Figure 1), and average losses are largest at the top. That tension is the interesting result. Each displaced worker falls from somewhere on the earnings ladder to the bottom of it, while wage and capital gains lift those who remain; because the displaced come disproportionately from higher-income households, the shock simultaneously thins the upper-middle of the distribution and swells its lower tail. When households are re-ranked on post-shock income, as the Gini requires, dispersion increases even though decile-average losses computed at fixed baseline ranks look progressive. A similar re-ranking mechanism is documented for Ireland by Doorley et al. (2026), and it distinguishes the AI case from earlier automation waves, in which exposure fell on mid- and low-wage routine work (Autor et al., 2003; Doorley et al., 2023).

6.1 Policy implications: what depends on who bears the shock

The central policy message follows from the incidence axis rather than from any single scenario. Inequality rises whichever family one assumes; what the incidence determines is whether the UK faces primarily a fiscal-resilience problem or a poverty-protection problem.

If the shock is top-loaded, as in the exposure-proportional family, the problem is chiefly fiscal. Displaced workers are drawn from occupations that pay well and are taxed heavily, so each displacement removes a large slug of income tax and National Insurance while adding comparatively little to means-tested expenditure. The Exchequer cost of the central shock is £18.4 billion under exposure-proportional incidence against £14.2 billion under uniform incidence—the same aggregate shock, roughly 30 per cent more expensive purely because of who bears it. Under top-loaded incidence the first-order policy question is revenue resilience: the UK tax base leans heavily on the taxation of labour income, and the workers most exposed are precisely those who contribute most to it. If AI additionally shifts the functional distribution of income from labour towards capital, the base that finances the cushioning I document would itself erode—a possibility that has motivated recent work on the taxation of capital in an AI-intensive economy (Bastani and Waldenström, 2024). The tax-composition exercise in Section 5 turns that possibility into a quantified scenario for the Treasury: because the effective labour-tax rate on displaced earnings is almost exactly the corporation-tax main rate, the net revenue shortfall is governed by the share of the lost wage bill that reappears as taxable corporate profit—full reappearance is roughly revenue-neutral, no reappearance leaves the whole labour-tax loss uncovered, and the harm is distributional rather than fiscal, since recycled dividends raise the Gini while reaching no household near the poverty line. The exercise is static accounting on top of the microsimulation, with corporation-tax incidence not modelled, but it gives base-erosion discussions a concrete UK arithmetic to anchor on.

If instead the shock lands flatter or on junior workers, the fiscal exposure moderates but the poverty exposure does not. Uniform incidence is the cheapest family I run, yet its poverty effect (+1.83 points before housing costs) is essentially the same as the exposure family’s (+1.87), and its Gini effect is the largest of the five families (+1.14 points). Junior-concentrated incidence costs slightly more than the exposure family (£19.4 billion) with a near-identical poverty effect. In these worlds the binding constraint is the adequacy of out-of-work support, and the relevant instruments are benefit levels, the contributory window, and—for the junior case—the entry prospects of young workers. Evidence that generative AI operates as a seniority-biased technology, compressing demand for junior and entry-level roles in exposed occupations (Hosseini and Lichtinger, 2026; Brynjolfsson et al., 2025a), with early UK corroboration in Klein Teeselink (2025), argues for close monitoring of youth labour-market entry. My junior-concentrated family can only re-weight who is displaced within a single year; the dynamic costs of a scarred entry cohort—foregone experience, delayed progression, persistent earnings losses—lie outside a static framework, so even that family understates the long-run burden on new entrants.

The fiscal-versus-poverty framing also now has a policy answer rather than only a diagnosis. The counterfactual reforms in Section 5 show that temporary Universal-Credit-side stabilisers—a circuit-breaker uplift to the standard allowance, or a benefit-cap suspension combined with a taper cut—deliver a point of after-housing-costs poverty reduction at roughly £10 billion, against £13–14 billion for an earnings-linked wage-insurance scheme, because wage insurance is earnings-proportional and much of its benefit accrues in the upper half of the baseline distribution. Wage insurance buys more total poverty reduction, but at disproportionate cost; a government facing a top-loaded shock and a strained revenue base gets more protection per pound by working through the existing means-tested architecture than by building an earnings-replacement scheme alongside it.

The expertise-compression family, my stylised stress test of the view that AI erodes the returns to accumulated expertise, is the worst of both worlds among the stylised families: the

most expensive (£24.5 billion) and the most poverty-raising (+2.12 points) of the four—and the measured family exceeds it on both margins (Table 2). If that view proves correct, neither fiscal buffers nor poverty protection alone suffices; both margins bind at once.

The shock also has a geography, and it does not track existing regional-policy maps. In the constituency analysis of Section 4.6, the largest proportionate income losses fall on Northern Ireland—the worst-hit region, at −2.3 per cent of aggregate household income—and Scotland, while Wales and the South East are mildest; the hardest-hit constituencies mix Northern Irish seats with inner-London ones. An AI labour shock is not a rerun of deindustrialisation: because exposure follows clerical and professional employment, it strikes administrative and professional centres wherever they sit, and a levelling-up strategy calibrated to the geography of manufacturing decline would miss much of it. These geographic results rest on imputed occupations on the enhanced 2023–24 dataset (Appendix A.8) and are indicative of relative, not absolute, local impacts.

Two implications hold across all families. First, because the modelled losses fall on identifiable occupational groups rather than diffusely across the workforce, retraining and lifelong-learning provision targeted at exposed occupations is the most direct lever available; the tax-benefit system can replace income but cannot restore occupational attachment. Second, since measured inequality rises under every incidence I examine, distributional monitoring should not wait for evidence on which incidence is materialising: the direction of the inequality effect is not in doubt within this framework, only its composition.

Taken together, the results suggest that the United Kingdom would enter a period of AI-driven labour-market disruption with a tax-benefit system that redistributes the losses effectively but does not, and cannot, prevent them. The system converts each market-income shock into a smaller disposable-income shock at a substantial, incidence-dependent fiscal cost, while the polarising character of displacement pushes measured inequality upwards regardless of who bears it. Policy attention is best directed at the margins the system does not reach: the occupational reallocation of exposed workers, the entry prospects of the young, and the robustness of a labour-taxation-heavy revenue base.

6.2 Limitations and further work

Several limitations qualify the results. I state them plainly, together with the measurement conventions a reader needs in order to compare my figures with official statistics.

Measurement conventions. All income results are on the HBAI cash disposable income concept at the household level, equivalised; poverty and inequality statistics share that concept. In computing the Gini I bottom-code negative equivalised incomes at zero. My baseline Gini is 0.309, below the official UK figure of roughly 0.34, because I use the plain Family Resources Survey without the SPI adjustment that corrects for under-coverage of top incomes; levels should therefore not be compared directly with official statistics, though changes—the objects of interest—are computed on a consistent basis throughout.

Monte Carlo conventions. Which results rest on which random draws matters for interpretation; the conventions are set out in Section 3.5 and Appendix A.1. Single-draw figures carry draw noise that the averaged figures do not, although the 20-draw runs indicate the aggregate magnitudes are stable across seeds.

Exposure resolution and crosswalk provenance. Occupational exposure is assigned at the 1-digit SOC level, because the Family Resources Survey records occupation only at the major-group level; within-group heterogeneity in AI exposure—substantial in the indices of [Felten et al. \(2021\)](#), [Webb \(2019\)](#) and [Eloundou et al. \(2024\)](#)—is averaged away. The exposure scores themselves reach the FRS through a reconstructed pipeline: the complementarity adjustment is rebuilt from published sources (recovering the published ranking up to a level offset), and the US occupational scores are carried to SOC2020 through a chained SOC2018–SOC2010–ISCO–SOC2020 crosswalk with ASHE employment weights available for 340 of 412 unit groups. A UK-native task-based index (GAISI) built from British Skills and Employment Survey data now exists ([Henseke et al., 2025](#)); adopting it, together with an imputation of detailed occupation from the Labour Force Survey, would remove the crosswalk and the resolution constraint at once, and is the first item on my agenda.

Survey weights. I use the uncalibrated FRS grossing weights rather than PolicyEngine’s enhanced dataset. The enhanced dataset improves aggregate fiscal accuracy through reweighting and household cloning, but cloning breaks the one-to-one correspondence between records and observed occupations on which the exposure join depends. Aggregate fiscal magnitudes should be read with the usual caveats attaching to unadjusted survey grossing.

Static displacement. Displaced workers are fully out of work in the simulated year—earnings, hours, employee pension contributions and salary sacrifice set to zero, no statutory pay, benefit entitlements recomputed under existing rules—while dual jobholders retain self-employment income. The duration structure of unemployment is not modelled, so the exhaustion of contributory benefits (new-style Jobseeker’s Allowance is payable for at most six months) cannot be expressed; over longer horizons I overstate the support available to displaced workers without means-tested entitlements.

The capital channel. The capital-income measure covers only interest and dividends, excluding rental income and returns inside private pensions, and household surveys under-represent top incomes and capital income in particular. The concentration of capital income in the top decile that drives this channel is, if anything, understated. Linked administrative or wealth-survey data (the Wealth and Assets Survey) would sharpen it considerably.

The self-employed. All shock scenarios exclude the self-employed. Exposure indices and the displacement literature underpinning my parameters are defined over employees, and self-employed earnings dynamics under AI adoption may include both displacement and new opportunities; I exclude the group rather than impose an arbitrary treatment. Since self-employment is a meaningful share of UK employment and is concentrated in some exposed occupations, this is a substantive omission.

Mechanisms and anchors. Where I attribute decile-level patterns to particular mechanisms—the thinness of savings buffers at the bottom, income pooling within couples—these attributions are hypotheses consistent with the data rather than results of a formal decomposition, and I flag them as such. The scenario anchors, likewise, are author-designed stress tests transported from headline estimates: the 1 per cent “low” cell is a sensitivity case rather than an employment estimate, the 7 per cent displacement and 2.6 per cent wage figures are conventions

converting task-exposure and productivity estimates into single-year shocks, and the paper-anchored scenarios are transpositions, not implementations, of the studies they draw on. As UK-specific evidence accumulates—early examples include Klein Teeselink (2025) and Rockall et al. (2025)—the scenarios can be re-anchored, and the gap between scenario analysis and conditional forecast will narrow. Extending the exposure measurement, restoring calibrated weights through an occupation-preserving reweighting, and introducing benefit-duration dynamics are the immediate priorities for further work.

Acknowledgements

The analysis uses the open-source PolicyEngine UK microsimulation model. Family Resources Survey microdata are accessed under the UK Data Service End User Licence and are not redistributed. All code and aggregate results that support the findings of this study are available at <https://github.com/PolicyEngine/uk-ai-study>. The author declares no conflict of interest.

References

- Acemoglu, D. (2025). The simple macroeconomics of AI. *Economic Policy*, 40(121), 13–58.
- Acemoglu, D. and Restrepo, P. (2018). The race between man and machine: Implications of technology for growth, factor shares, and employment. *American Economic Review*, 108(6), 1488–1542.
- Acemoglu, D. and Restrepo, P. (2022). Tasks, automation, and the rise in US wage inequality. *Econometrica*, 90(5), 1973–2016.
- Acemoglu, D., Autor, D., Hazell, J. and Restrepo, P. (2022). Artificial intelligence and jobs: Evidence from online vacancies. *Journal of Labor Economics*, 40(S1), S293–S340.
- Agrawal, A., Gans, J. and Goldfarb, A. (2018). *Prediction Machines: The Simple Economics of Artificial Intelligence*. Harvard Business Review Press, Boston, MA.
- Autor, D. H., Levy, F. and Murnane, R. J. (2003). The skill content of recent technological change: An empirical exploration. *The Quarterly Journal of Economics*, 118(4), 1279–1333.
- Autor, D. (2024). Applying AI to rebuild middle class jobs. NBER Working Paper No. 32140, National Bureau of Economic Research, Cambridge, MA.
- Bastani, S. and Waldenström, D. (2024). AI, automation and taxation. CESifo Working Paper No. 11084 / IZA Policy Paper No. 212, Munich and Bonn.
- Briggs, J. and Kodnani, D. (2023). The potentially large effects of artificial intelligence on economic growth. Goldman Sachs Global Economics Analyst, Goldman Sachs.
- Brynjolfsson, E., Chandar, B. and Chen, R. (2025a). Canaries in the coal mine? Six facts about the recent employment effects of artificial intelligence. Working paper, Stanford Digital Economy Lab.
- Brynjolfsson, E., Li, D. and Raymond, L. (2025b). Generative AI at work. *The Quarterly Journal of Economics*, 140(2), 889–942.

- Cazzaniga, M., Jaumotte, F., Li, L., Melina, G., Panton, A. J., Pizzinelli, C., Rockall, E. and Tavares, M. M. (2024). Gen-AI: Artificial intelligence and the future of work. IMF Staff Discussion Note SDN/2024/001, International Monetary Fund, Washington, DC.
- Doorley, K., Gromadzki, J., Lewandowski, P., Tuda, D. and Van Kerm, P. (2023). Automation and income inequality in Europe. IZA Discussion Paper No. 16499, Institute of Labor Economics, Bonn.
- Doorley, K., O'Connor, S., O'Shea, R. and Tuda, D. (2026). Artificial intelligence and income inequality in Ireland. Jointly-published Report 16, Economic and Social Research Institute and Department of Finance, Dublin.
- Department for Science, Innovation and Technology and Department for Education (2023). The impact of AI on UK jobs and training. UK Government, London.
- Eloundou, T., Manning, S., Mishkin, P. and Rock, D. (2024). GPTs are GPTs: Labor market impact potential of LLMs. *Science*, 384(6702), 1306–1308.
- Felten, E., Raj, M. and Seamans, R. (2021). Occupational, industry, and geographic exposure to artificial intelligence: A novel dataset and its potential uses. *Strategic Management Journal*, 42(12), 2195–2217.
- Frey, C. B. and Osborne, M. A. (2017). The future of employment: How susceptible are jobs to computerisation? *Technological Forecasting and Social Change*, 114, 254–280.
- Goos, M. and Manning, A. (2007). Lousy and lovely jobs: The rising polarization of work in Britain. *The Review of Economics and Statistics*, 89(1), 118–133.
- Goos, M., Manning, A. and Salomons, A. (2014). Explaining job polarization: Routine-biased technological change and offshoring. *American Economic Review*, 104(8), 2509–2526.
- Henseke, G. et al. (2025). How exposed are UK jobs to generative AI? The GAISI index. arXiv:2507.22748.
- Hosseini, S. and Lichtinger, G. (2026). Generative AI as seniority-biased technological change: Evidence from U.S. résumé and job posting data. SSRN Working Paper No. 5425555.
- Humlum, A. and Vestergaard, E. (2025). Large language models, small labor market effects. Working paper, University of Chicago Becker Friedman Institute.
- International Labour Organization (2025). Generative AI and jobs: A 2025 update. ILO, Geneva.
- Klein Teeselink, B. (2025). Generative AI and labor market outcomes: Evidence from the UK. SSRN Working Paper No. 5516798.
- Korinek, A. and Stiglitz, J. E. (2019). Artificial intelligence and its implications for income distribution and unemployment. In Agrawal, A., Gans, J. and Goldfarb, A. (eds), *The Economics of Artificial Intelligence: An Agenda*, University of Chicago Press, Chicago, 349–390.
- Moll, B., Rachel, L. and Restrepo, P. (2022). Uneven growth: Automation's impact on income and wealth inequality. *Econometrica*, 90(6), 2645–2683.

- Noy, S. and Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187–192.
- Office for Budget Responsibility (2025). Fiscal risks and sustainability report. OBR, London.
- Office for Budget Responsibility (2026). Fiscal risks and sustainability report, July 2026. OBR, London.
- Pizzinelli, C., Panton, A. J., Tavares, M. M., Cazzaniga, M. and Li, L. (2023). Labor market exposure to AI: Cross-country differences and distributional implications. IMF Working Paper WP/23/216, International Monetary Fund, Washington, DC.
- Rockall, E., Pizzinelli, C. and Tavares, M. M. (2025). Artificial intelligence and wage inequality in the United Kingdom. IMF Working Paper, International Monetary Fund, Washington, DC.
- Tolan, S., Pesole, A., Martínez-Plumed, F., Fernández-Macías, E., Hernández-Orallo, J. and Gómez, E. (2021). Measuring the occupational impact of AI: Tasks, cognitive abilities and AI benchmarks. *Journal of Artificial Intelligence Research*, 71, 191–236.
- Webb, M. (2019). The impact of artificial intelligence on the labor market. Working paper, Stanford University.
- Williamson, S., et al. (2025). Applying the complementarity-adjusted AI occupational exposure index to the United Kingdom and Ireland. *Journal for Labour Market Research*, 59, 30.

A Additional tables and figures

A.1 Seed and draw conventions

Displacement is a random draw of which exposed workers lose employment, subject to occupation-group quotas proportional to employment times mean exposure. Throughout the paper, three conventions apply. The scenario grid (Figures 5–7, and their appendix variants below), the decomposition in Figure 4, the preset scenarios and the four incidence families are single draws at seed 0. The decile transition and wage-gain profiles (Figures 1 and 2) are means over 50 independent draws. The robustness summary and the confidence intervals in Figure 19 are 20-draw Monte Carlo runs, with $\pm$ values reported as standard deviations across draws. Within each draw, ordering keys are uniform within occupation group, so a record’s inclusion probability does not depend on its grossing weight.

A.2 Job loss by occupation group

Table 5 reports the central-scenario displacement rate for each SOC2020 major group. The allocation rule concentrates losses in administrative and secretarial work (10.5 per cent of the group’s employment), managerial and professional occupations (around 9.5 per cent) and associate professional roles, while elementary occupations—the least-exposed group, whose normalised exposure score is zero—experience no displacement at all. Clerical and professional work bears the shock; manual and elementary work is insulated.

Table 5: Central-scenario job loss by SOC2020 major group

Major group	C-AIOE	Employment (m)	Job loss (%)
1 Managers, directors and senior officials	0.62	2.45	9.6
2 Professional occupations	0.60	6.46	9.5
3 Associate professional occupations	0.47	3.13	8.5
4 Administrative and secretarial occupations	0.74	2.35	10.5
5 Skilled trades occupations	−0.33	1.56	2.8
6 Caring, leisure and other service occupations	−0.14	2.32	3.9
7 Sales and customer service occupations	0.27	1.22	7.2
8 Process, plant and machine operatives	−0.54	1.41	1.3
9 Elementary occupations	−0.73	2.05	0.0

C-AIOE is the standardised complementarity-adjusted exposure score before the min-zero normalisation used in the allocation rule. Employment and job-loss rates are survey-weighted; single draw, seed 0.

A.3 Where the shock lands: baseline distributions

Figures 12–16 show the baseline household-level decile distributions of the income components through which the three channels operate (aggregate £billion per year by decile of equivalised household disposable income). Market income excluding capital (£1,313 billion in total) and income tax and National Insurance (£312 billion) are heavily concentrated in the top deciles—which is why displacement of high earners is fiscally expensive. Benefits (£320 billion) are spread across the lower and middle deciles, peaking in deciles 3 to 5. Capital income (£23 billion of interest and dividends) is the most concentrated series of all: the top decile holds more than a third of the total.




Figure 12: Baseline aggregate household market income (less capital income) by decile.




Figure 13: Baseline aggregate household capital income (interest and dividends) by decile.




Figure 14: Baseline aggregate household benefits by decile.




Figure 15: Baseline aggregate household income tax and National Insurance by decile.




Figure 16: Baseline aggregate household disposable income (HBAI concept) by decile.

A.4 Exposure incidence before any simulation

Figures 17 and 18 show, without any shock, how workers in each income decile distribute across employment-weighted tertiles of AI exposure and of complementarity. The share of workers in the top exposure tertile rises steadily with household income, while complementarity shows no comparable gradient—the descriptive facts that generate the simulated incidence in Section 4.




Figure 17: Distribution of employed persons across AI-exposure tertiles, by income decile.




Figure 18: Distribution of employed persons across complementarity tertiles, by income decile.

A.5 Monte Carlo uncertainty on the decile profile

Figure 19 reports the central-scenario disposable-income change by decile as the mean over twenty independent displacement draws with 95 per cent confidence intervals. The gradient runs from -0.5 per cent in the bottom decile to -4.1 per cent in the top. It is not quite

monotone: deciles 2 and 3 show a small reversal (-0.88 against -0.82 per cent), well inside the draw noise. The overall upward slope of the loss profile survives the uncertainty comfortably.




Figure 19: Disposable income change by decile, central scenario: mean of 20 displacement draws with 95 per cent confidence intervals.

A.6 Uniform-shock and alternative-index comparisons by decile

Figure 20 compares the central scenario against the same aggregate shocks allocated uniformly across occupations, decile by decile; Figure 21 repeats the central scenario with the [Eloundou et al. \(2024\)](#) GPT-exposure index in place of the C-AIOE. The exposure-based allocation shifts losses up the distribution relative to a uniform shock (top-decile losses of 3.9 against 3.0 per cent; bottom-decile losses of 0.5 against 0.9 per cent), and the choice of exposure index leaves the decile profile essentially unchanged.

A.7 Full exposure-index sensitivity

Figure 21 compared decile profiles under two indices; here I rerun the entire central scenario under five: the C-AIOE (my central measure), the unadjusted [Felten et al. \(2021\)](#) AIOE, the [Eloundou et al. \(2024\)](#) GPT-exposure beta, and the two indices from the UK government’s DSIT exposure assessment (all-AI and large-language-model variants). The comparison matters because the published literature disagrees sharply on *levels*: estimates of the share of employment “highly exposed” to AI range from a few per cent to well over half depending on the index and threshold chosen, so any result that survived under only one index would be fragile. My quota mechanism holds the aggregate displacement rate at 7 per cent, so the indices differ only in *where* the same aggregate shock lands across occupation groups.

Table 6 and Figure 22 report the outcomes. The exchequer cost spans £17.6–18.6 billion (the largest deviation from the C-AIOE figure is -4.3 per cent, under the DSIT LLM index), the Gini change spans $+1.05$ to $+1.10$ points, and the before-housing-costs poverty change spans $+1.85$ to $+1.90$ points. The decile-10 transition share exceeds the decile-1 share by a factor of 8.9 to 11.3 across indices. All three qualitative conclusions of the paper—the Gini rises, the




Figure 20: Disposable income change by decile: exposure-allocated versus uniform shock, central parameters.




Figure 21: Disposable income change by decile under the C-AIOE and the Eloundou et al. exposure index.

displacement-transition gradient rises with income, and the fiscal cost is within 20 per cent of the central estimate—hold under every index, essentially because all five rank SOC major groups similarly: clerical and professional work is most exposed under each, and the quota fixes the aggregate.

Table 6: Central scenario under five exposure indices

Index	Cost (£bn)	Δ Pov. BHC (pp)	Δ Gini (pp)	D1 share (%)	D10 share (%)
C-AIOE (central)	18.4	1.87	1.10	0.53	4.67
Felten AIOE	18.6	1.86	1.09	0.53	4.70
Eloundou beta	18.4	1.90	1.09	0.47	4.74
DSIT AIOE	18.6	1.86	1.05	0.43	4.86
DSIT LLM	17.6	1.85	1.09	0.51	4.68

Single draw, seed 0, 2026. D1 and D10 shares are the percentages of persons in the bottom and top baseline income deciles whose household is pushed down at least one decile by displacement. Aggregate displacement is held at 7 per cent by the quota mechanism under every index.




Figure 22: Headline outcomes of the central scenario under five exposure indices.

A.8 Data appendix: constituency-level occupation imputation

The constituency results use PolicyEngine’s 650-constituency weight matrix, which indexes the enhanced FRS 2023–24 dataset (53,566 employee records); on that dataset the serial-number join to observed SOC major group is invalid, so occupation had to be imputed. I drew each enhanced record’s SOC major group from the survey-weighted occupation distribution of the plain FRS 2024–25 (12,800 employees, 99.6 per cent with observed SOC) within cells defined by age band $\times$ gender $\times$ region $\times$ individual-earnings decile, at seed 0. On a weighted basis 99.4 per cent of employees were imputed from the full four-way cell, 0.6 per cent from an age–gender–decile fallback and 0.03 per cent from the marginal distribution. As stressed in the main text, occupation is therefore imputed rather than observed at the constituency level: cross-constituency variation reflects demographic and earnings composition (plus one seed-0 draw), not directly measured local occupation mixes, and the enhanced-dataset results are not directly comparable to the plain-FRS poverty numbers elsewhere in the paper.

A.9 The scenario grid by decile, and without the capital shock

Figure 23 facets the full 66-cell scenario grid by baseline income decile; cells report the change in mean household disposable income (HBAI concept), which across the grid as a whole ranges from +3.0 to −5.2 per cent. Every decile’s panel shifts towards losses as displacement rises, but the top deciles darken much faster: in the harshest cell (10 per cent displacement, no wage gain) the bottom decile loses 1.9 per cent of disposable income while the top decile loses 8.2 per cent. Figure 24 removes the capital-return shock. Average disposable incomes fall relative to the capital-on grid, confirming that the capital channel raises average incomes while concentrating its gains at the top; the Gini result, however, is unchanged in kind: without the capital shock the Gini of equalised disposable income still rises in all 66 cells, by 0.00 to 1.50 percentage points. The inequality result does not depend on the capital channel.




Figure 23: Change in household disposable income across the scenario grid, faceted by baseline income decile. Single draw, seed 0.




Figure 24: Average change in household disposable income across the scenario grid with the capital-return shock switched off. Single draw, seed 0.